

Naive Bayes classifiers are built on Bayesian classification methods. These rely on Bayes's theorem, which is an equation describing the relationship of conditional probabilities of statistical quantities.

Create a Model using Naive Bayes classifiers to predict whether a passenger on the titanic would have been survived or not.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
%matplotlib inline
```

```
titanic = pd.read_csv("/content/train.csv")
titanic.head()
```

|   | PassengerId | Survived | Pclass | Name  | Sex    | Age  | SibSp | Parch | Ticket    |
|---|-------------|----------|--------|---|--------|------|-------|-------|-----------|
| 0 | 1           | 0        | 3      | Braund,<br>Mr. Owen<br>Harris                                 | male   | 22.0 | 1     | 0     | A/5 21171 |
| 1 | 2           | 1        | 1      | Cumings,<br>Mrs. John<br>Bradley<br>(Florence<br>Briggs<br>Th | female | 38.0 | 1     | 0     | PC 17599  |

```
print("Shape of dataset:", titanic.shape) # shape of dataset
```

Shape of dataset: (891, 12)

```
print("Columns present in dataset:\n", titanic.columns) # columns present in dataset
```

```
Columns present in dataset:
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
       'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
      dtype='object')
```

```
titanic.isnull().sum() # check total null values inside data
```

```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age           177
SibSp           0
Parch           0
Ticket          0
```

```
Fare      0
Cabin    687
Embarked  2
dtype: int64
```

```
# fill values of age column
```

```
titanic.fillna(titanic.mean(), inplace = True)
titanic.isnull().sum()
```

```
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             0
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin            687
Embarked         2
dtype: int64
```

```
# fill values of Embarked column
```

```
titanic["Embarked"].fillna("S", inplace = True)
titanic.isnull().sum()
```

```
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             0
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin            687
Embarked         0
dtype: int64
```

```
# drop Cabin column because it has lot of null values. 687/891
```

```
drop_cabin = titanic.isnull().sum()[titanic.isnull().sum() > (50/100 * titanic.shape[0])]
drop_cabin
```

```
Cabin      687
dtype: int64
```

```
drop_cabin.index
```

```
Index(['Cabin'], dtype='object')
```

```
titanic.drop(drop_cabin.index, axis = 1, inplace = True)
titanic.isnull().sum()
```

```
PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Embarked       0
dtype: int64
```

```
titanic.describe()
```

|              | PassengerId | Survived   | Pclass     | Age        | SibSp      | Parch      | Fare       |
|--------------|-------------|------------|------------|------------|------------|------------|------------|
| <b>count</b> | 891.000000  | 891.000000 | 891.000000 | 891.000000 | 891.000000 | 891.000000 | 891.000000 |
| <b>mean</b>  | 446.000000  | 0.383838   | 2.308642   | 29.699118  | 0.523008   | 0.381594   | 32.204208  |
| <b>std</b>   | 257.353842  | 0.486592   | 0.836071   | 13.002015  | 1.102743   | 0.806057   | 49.693429  |
| <b>min</b>   | 1.000000    | 0.000000   | 1.000000   | 0.420000   | 0.000000   | 0.000000   | 0.000000   |
| <b>25%</b>   | 223.500000  | 0.000000   | 2.000000   | 22.000000  | 0.000000   | 0.000000   | 7.910400   |
| <b>50%</b>   | 446.000000  | 0.000000   | 3.000000   | 29.699118  | 0.000000   | 0.000000   | 14.454200  |
| <b>75%</b>   | 668.500000  | 1.000000   | 3.000000   | 35.000000  | 1.000000   | 0.000000   | 31.000000  |
| <b>max</b>   | 891.000000  | 1.000000   | 3.000000   | 80.000000  | 8.000000   | 6.000000   | 512.329200 |

```
titanic.corr()
```

| PassengerId | Survived | Pclass | Age | SibSp | Parch | Fare |
|-------------|----------|--------|-----|-------|-------|------|
|-------------|----------|--------|-----|-------|-------|------|

```
titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age             891 non-null    float64
6   SibSp           891 non-null    int64
7   Parch           891 non-null    int64
8   Ticket          891 non-null    object
9   Fare            891 non-null    float64
10  Embarked        891 non-null    object
dtypes: float64(2), int64(5), object(4)
memory usage: 76.7+ KB
```

```
# create a new column Family size by adding SibSp and Parch
```

```
titanic["FamilySize"] = titanic["SibSp"] + titanic["Parch"]
titanic.head()
```

|   | PassengerId | Survived | Pclass | Name  | Sex    | Age  | SibSp | Parch | Ticket    | Fare   |
|---|-------------|----------|--------|---|--------|------|-------|-------|-----------|--------|
| 0 | 1           | 0        | 3      | Braund,<br>Mr. Owen<br>Harris                                 | male   | 22.0 | 1     | 0     | A/5 21171 | 7.250  |
| 1 | 2           | 1        | 1      | Cumings,<br>Mrs. John<br>Bradley<br>(Florence<br>Briggs<br>Th | female | 38.0 | 1     | 0     | PC 17599  | 71.283 |

```
# drop SibSp and Parch because we create new column FamilySize instaed of them
```

```
titanic.drop(["SibSp", "Parch"], axis = 1, inplace = True)
titanic.head()
```

| PassengerId | Survived | Pclass | Name | Sex | Age | Ticket | Fare | Embarked | F |
|-------------|----------|--------|------|-----|-----|--------|------|----------|---|
|-------------|----------|--------|------|-----|-----|--------|------|----------|---|

Braund

```
titanic.corr()
```

|             | PassengerId | Survived  | Pclass    | Age       | Fare      | FamilySize |
|-------------|-------------|-----------|-----------|-----------|-----------|------------|
| PassengerId | 1.000000    | -0.005007 | -0.035144 | 0.033207  | 0.012658  | -0.040143  |
| Survived    | -0.005007   | 1.000000  | -0.338481 | -0.069809 | 0.257307  | 0.016639   |
| Pclass      | -0.035144   | -0.338481 | 1.000000  | -0.331339 | -0.549500 | 0.065997   |
| Age         | 0.033207    | -0.069809 | -0.331339 | 1.000000  | 0.091566  | -0.248512  |
| Fare        | 0.012658    | 0.257307  | -0.549500 | 0.091566  | 1.000000  | 0.217138   |
| FamilySize  | -0.040143   | 0.016639  | 0.065997  | -0.248512 | 0.217138  | 1.000000   |

```
# filtered alone persons/passengers
```

```
titanic["Alone"] = [0 if titanic["FamilySize"][i] > 0 else 1 for i in titanic.index]
titanic.head()
```

|   | PassengerId | Survived | Pclass | Name  | Sex    | Age  | Ticket    | Fare    | Embarked | F |
|---|-------------|----------|--------|---|--------|------|-----------|---------|----------|---|
| 0 | 1           | 0        | 3      | Braund,<br>Mr. Owen<br>Harris                                 | male   | 22.0 | A/5 21171 | 7.2500  |          | S |
| 1 | 2           | 1        | 1      | Cumings,<br>Mrs. John<br>Bradley<br>(Florence<br>Briggs<br>Th | female | 38.0 | PC 17599  | 71.2833 |          | C |


```
titanic.corr()
```

|             | PassengerId | Survived  | Pclass    | Age       | Fare      | FamilySize | Alone     |
|-------------|-------------|-----------|-----------|-----------|-----------|------------|-----------|
| PassengerId | 1.000000    | -0.005007 | -0.035144 | 0.033207  | 0.012658  | -0.040143  | 0.057462  |
| Survived    | -0.005007   | 1.000000  | -0.338481 | -0.069809 | 0.257307  | 0.016639   | -0.203367 |
| Pclass      | -0.035144   | -0.338481 | 1.000000  | -0.331339 | -0.549500 | 0.065997   | 0.135207  |
| Age         | 0.033207    | -0.069809 | -0.331339 | 1.000000  | 0.091566  | -0.248512  | 0.179775  |
| Fare        | 0.012658    | 0.257307  | -0.549500 | 0.091566  | 1.000000  | 0.217138   | -0.271832 |
| FamilySize  | -0.040143   | 0.016639  | 0.065997  | -0.248512 | 0.217138  | 1.000000   | -0.690922 |
| Alone       | 0.057462    | -0.203367 | 0.135207  | 0.179775  | -0.271832 | -0.690922  | 1.000000  |

Filtered out survived ratio according to conditions and visualize them


```
# sex ratio of passengers
```

```
sb.countplot(x = "Sex", data = titanic);
```



```
# age distribution
```

```
plt.hist(x = titanic["Age"], bins = 20);
```




```
# passenger class
```

```
x = titanic["Pclass"].value_counts()
```

```
plt.pie(x, labels = x.index, startangle = 90, counterclock = False);
```


```
plt.legend()
```

<matplotlib.legend.Legend at 0x7fc5c7899210>




```
#Embarked
y = titanic["Embarked"].value_counts()
plt.pie(y, labels = y.index, startangle = 90, counterclock = True);
plt.legend()
```

<matplotlib.legend.Legend at 0x7fc5c77fec90>




```
# survive rate of alone person according to their sex
sb.countplot(x = "Alone", hue = "Sex", data = titanic);
```




```
# survive rate of family
```

```
sb.countplot(x = "FamilySize", data = titanic);
```



```
# total survived passengers
```

```
sb.countplot(x = "Survived", data = titanic);
```




```
# survived ratio according to sex
```

```
sb.countplot(x = "Survived", hue = "Sex", data = titanic);
```






```
# according to pclass
sb.countplot(x = "Survived", hue = "Pclass", data = titanic);
```




```
# according to embarked
sb.countplot(x = "Survived", hue = "Embarked", data = titanic);
```



```
# according to sex and passenger class
sb.barplot("Sex", "Pclass", hue = "Survived", data = titanic);
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {\"Pclass\", \"Sex\", \"Age\", \"Fare\", \"Embarked\", \"FamilySize\", \"Alone\"}. Future versions will require their presence.



## Label Encoding for Sex and Embarked

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
titanic["Sex"] = le.fit_transform(titanic["Sex"])
titanic["Embarked"] = le.fit_transform(titanic["Embarked"])
print("Encoded values for Sex:", titanic["Sex"].unique())
print("Encoded values for Embarked:", titanic["Embarked"].unique())
```

```
Encoded values for Sex: [1 0]
Encoded values for Embarked: [2 0 1]
```

```
titanic.head()
```

|   | PassengerId | Survived | Pclass | Name  | Sex | Age  | Ticket    | Fare    | Embarked | FamilySize |
|---|-------------|----------|--------|---|-----|------|-----------|---------|----------|------------|
| 0 | 1           | 0        | 3      | Braund, Mr. Owen Harris                         | 1   | 22.0 | A/5 21171 | 7.2500  | 2        | 1          |
| 1 | 2           | 1        | 1      | Cumings, Mrs. John Bradley (Florence Briggs Th) | 0   | 38.0 | PC 17599  | 71.2833 | 0        | 1          |

## Features and Target

```
features = titanic[["Pclass", "Sex", "Age", "Fare", "Embarked", "FamilySize", "Alone"]]
target = titanic["Survived"]
```

## Divide data for training and testing

```
from sklearn.model_selection import train_test_split

xtrain, xtest, ytrain, ytest = train_test_split(features, target, test_size = 0.3, random_state = 42)
print("Shape of xtrain:", xtrain.shape)
```

```
print("Shape of ytrain:", ytrain.shape)
print("Shape of xtest:", xtest.shape)
print("Shape of ytest:", ytest.shape)
```

```
Shape of xtrain: (623, 7)
Shape of ytrain: (623,)
Shape of xtest: (268, 7)
Shape of ytest: (268,)
```

## Create a model and train the data

```
from sklearn.naive_bayes import GaussianNB
```

```
gnb = GaussianNB()
gnb.fit(xtrain, ytrain)
```

```
GaussianNB(priors=None, var_smoothing=1e-09)
```

## Test the testing data and make prediction

```
ypred = gnb.predict(xtest)
print("Prediction made by model:\n", ypred)
```

```
Prediction made by model:
[0 0 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1
 1 0 1 0 0 0 0 0 0 1 0 1 1 1 0 1 1 0 0 1 1 1 0 0 1 0 0 0 0 1 0 0 0
 0 0 0 0 1 0 0 0 0 0 0 1 1 0 1 1 0 0 0 0 1 0 1 1 0 0 0 1 0 1 0 0 0 1 1 0 0
 0 1 1 0 1 0 0 0 0 1 0 1 0 1 1 1 1 0 0 1 1 1 1 0 0 1 0 1 1 0 1 0 1 1 1 1
 0 0 1 1 1 0 0 1 0 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 1 0 0 0 1 0 1 0 0 1 0 1
 0 1 0 1 0 0 1 0 1 1 1 0 1 1 0 0 0 0 1 1 0 0 0 0 1 0 1 0 0 0 0 0 0 1 1 1 0
 0 0 1 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 1 0 1 0 0 0 0 1 1 1 1 0 0 0 1 1
 0 0 0 1 0 0 0 1 0]
```

## Confusion Matrix and Accuracy

```
from sklearn.metrics import confusion_matrix, accuracy_score
```

```
matrix = confusion_matrix(ytest, ypred)
print("Confusion Matrix of a model:\n", matrix)
```

```
Confusion Matrix of a model:
[[154  24]
 [ 20  70]]
```

```
accuracy = accuracy_score(ytest, ypred)
print("Accuracy of mdoel: {}".format(accuracy*100))
```

```
Accuracy of mdoel: 83.5820895522388%
```

